

Handin 3 - Semi-supervised deep learning

Rasmus Freund - 201700273

December 4, 2023

Short answer questions

- Where are we ensuring that the finetuning does not affect the feature extractor?
 - The `Net` class has been created with separate components for feature extraction and classification; which has additionally been split into two layers: `pretrainer` and `generalizer`. When finetuning (`optim.SGD(network.generalizer.parameters(),lr=0.01)`) we only update the `generalizer` parameters, not those of the feature extractor.
- How does the code work that gets s samples per class for the pre-training dataset
 - The `get_subsampled_dataset` function iterates over each class in the dataset. For each class, it finds the indices where the target label matches the current class and randomly selects k indices from that list which are then appended to the subset.
- What is the `forward_call` parameter responsible for in the `train()` method (located in `network_training.py`)?
 - It performs a forward pass through the network, calculating the loss.
- Describe how the `augment()` method works (located `augmentations.py`).
 - The function iterates through the batch of images, randomly selects another image from the batch (`merge_indices[i]=np.random.choice(...)`). The current image (i) is excluded from the pool, so we're never getting an augment where an image has been augmented with itself. Augmentation is then applied (either mixup or collage), resulting in a new image and an interpolation value.
- If we pre-train and finetune on the same dataset, is there any reason to do the finetuning step?
 - Typically, no. However, there might be specific cases where this could be beneficial. One example could be transfer learning where we could pre-train the model as a simple classification network. After that, the fine-tuning head could then be utilized for more specific tasks.

Predictions

- Will the collage and mixup data augmentations help achieve higher finetune accuracies? Which do you expect will be more effective?
 - I doubt that data augmentation is going to improve accuracy in the case of these data sets. Changing the information in the images through augmentation is also going to change the true "meaning" of those images, since we're dealing with numbers and letters. Of the two types of augmentation, I suspect that mixup might perform better, since the entire structure of each number/letter will be kept in the augmented data.
- What relationship do you expect between the number of samples in the pre-training dataset and the finetuning accuracy? Does this change with data augmentations?
 - Generally, more samples in the pre-training should improve the accuracy of the fine-tuning head. However, given that the data is relatively simple and present clear patterns, a plateau is likely to be reached fairly quickly in terms of amount of training samples being used. Beyond this plateau, increasing the amount of samples will probably not lead to significant increases in accuracy. While data augmentation is, under most circumstances, a good way to increase accuracy when training on sparse data, I suspect it won't have a great impact in this network, mainly due to the reasons listed in the answer to the previous question.